Inside Outside Recursive Neural Network: A Unified Framework for Compositional Semantics and Meaning in Context

Phong Le, Remko Scha and **Willem Zuidema**

Abstract

Although compositional semantics and meaning in context have a strong relation, it is surprising that there are no attempts tackling both. In this paper, we propose a novel framework, Inside Outside Recursive Neural Network, which is able to compute phrase representations and context representations, thus solving the both challenges. Center to our framework is a new hypothesis that, according to our experimental results, is promising to lead to efficient unsupervised learning schemes for neural compositional semantics.

1 Introduction

The reader should not have any problems to understand *this* sentence although (s)he has never seen it. This is evidence that natural languages are compositional. In order to capture this phenomenon, compositional semantics which relies on the principle of compositionality *"The meaning of a whole is a function of the meanings of the parts and of the way they are syntactically combined"* (Partee, 1995) was introduced. Research on compositional semantics therefore focuses on two problems (1) how to learn word representations, and (2) how to learn compositionality functions.

Phrase/word meaning in context, on the other hand, is about how meaning of a phrase/word is changed by effect of it contexts. For instance, at the word level, the word 'bank' in the two following constituents has different meanings

1. along the east *bank* of the river

2. some big institutions, including *banks*

At the phrase level, depending on its context, a phrase should be understood literally or figuratively.

It is not difficult to realize that compositional semantics and meaning in context have a strong relation. To meaning in context, compositionality is helpful to compute context representations. In addition, if the target is a phrase, compositionality is essential for computing its meaning. To compositional semantics, context can be used to disambiguate word senses, thus lead to more reliable composition. However, it is surprising that there are no attempts tackling both problems[1].

In this paper, we propose a new framework, namely Inside Outside Recursive Neural Network (IORNN). IORNN is able to compute phrase representations as well as context representations, thus tackles the both challenges at the same time in a unified framework.

2 Background

2.1 Compositional Semantics

Formal semantics, which uses formal languages to represent constituents, is the first attempt and well fits the principle of compositionality (Montague, 1970). It firstly assumes that words are represented by lambda expressions, e.g. "John" :- $\lambda x.john(x)$, "walks" :- $\lambda P \lambda y.walks(y) \wedge P(y)$. Then, using the lambda beta reduction rule as the compositionality function, it easily computes the meaning of a constituent ("John walks" :- $\lambda y.john(y) \wedge walks(y)$). This approach, although being sound with human beings and having a simple but powerful compositionality function, is very challenging to computers since automatically learning word

representations is a difficult task (Le and Zuidema, 2012). In addition, the fact that formal semantics only supports the truth values (True and False) and ignores lexical semantics restricts it to a short list of applications.

At the other extreme, distributional semantics, based on the distributional hypothesis (Lenci, 2008) *"The degree of semantic similarity between two linguistic expressions A and B is a function of the similarity of the linguistic contexts in which A and B can appear"*, is widely used for learning word meanings, which are represented in vector spaces. That is because it has strong support from not only linguistics but also cognitive science (Lenci, 2008). In addition, word representations can be learnt from unannotated data, which are redundant and free on the Internet, thanks to the flourish of unsupervised learning techniques, such as latent semantic analysis (Landauer et al., 1998), neural network language model (Collobert et al., 2011; Huang et al., 2012), Brown clustering algorithm (Brown et al., 1992), and spectral learning (Dhillon et al., 2012). And after all, due to the fact that it supports semantic similarity, distributional semantics has a wide range of applications: information retrieval, sentiment analysis, machine translation, etc. The distributional hypothesis, however, is not able to apply to phrasal semantics because of the sparsity of data: the number of phrases is infinite whereas available data is finite.

The most visible approach to tackle compositional semantics problem is to combine formal semantics and distributional semantics: the former for compositionality, the latter for word representations (Garrette et al., 2012). It turns out this is not trivial since the two kinds of semantics have very different forms of representations so that how to combine these forms is itself a difficult problem.

Another approach, which has been intensively studied recently, is distributional compositional semantics: if $\overrightarrow{a}, \overrightarrow{b}$ are vectors representing the meanings of the two items a, b, then the meaning vector $\overrightarrow{ab}$ of their constituent ab, yielded by the grammar rule R, is computed by (Mitchell and Lapata, 2008)

$$\overrightarrow{ab} = f(\overrightarrow{a}, \overrightarrow{b}, R, K) \qquad (1)$$

where f is a compositionality function and K is background knowledge. Many approaches were proposed to learn compositionality functions f.

The most simple one is to use vector addition and multiplication as compositionality functions (Mitchell and Lapata, 2008).

Socher and colleagues propose two neural network frameworks: recursive auto encoder (RAE) (Socher et al., 2011a) for unsupervised learning, and recursive neural network (RNN) for supervised learning with task-based training signal (Socher et al., 2010) (e.g., for sentiment analysis, the training signal is the sentiments given by voters). The key idea of the RAE framework is: a compositionality function is a compression function, such that an input is able to be recovered from the output by a decompression function.

Baroni et al., 2012), Grefenstette et al., 2013) and others attempt the challenge in a different way. They use tensors to represent functor words (i.e., verbs, adjectives, etc.), linear maps as compositionality functions, and use contexts of phrases (in a similar way as in distributional lexical semantics) for estimating tensors' elements and functions' parameters.

Our work has some common points with the two approaches above. Similarly to the work of Socher and colleagues, IORNN uses recursive network architecture to construct phrase representations, but different in that it also constructs context representations and is trained in a very different manner. The training has the same key idea with the works of Baroni et al., 2012), Grefenstette et al., 2013): context is used as training signal. However, we only use contexts of words and thus avoid the sparsity of data.

2.2 Meaning in Context

[TODO...]

3 Research Questions

First of all, we ask ourselves "Which evidence is strong enough to be used for learning compositional semantics?" To answer this, we rely on the following observation: a human being can guess the meaning of an unknown word by making use of the meaning of its context. In other words, he computes (in his brain) the meaning of the context and then uses it to predict the meaning of the target unknown word (by setting some constraints to narrow down a list of possible meanings). Hence, if he correctly predicts the meaning of the unknown word, we can, at some degree of belief, say that he comprehends the meaning of the context. This

idea is then captured in the following hypothesis

> *Hypothesis 1:* The agreement between words and contexts provides evidence for unsupervised compositional semantics learning.

Then, there are three questions need to be answered in order to implement the hypothesis

1. How to construct phrase representations?

2. How to construct context representations?

3. How to use the agreement between words and their contexts to learn compositionality functions?

We present our answers in the next section.

4 Methodology

4.1 Recursive Neural Network (RNN)

Socher et al., 2010) answer the first question "How to construct phrase representations?" by Recursive Neural Network (RNN) architecture. In order to see how RNN works, let's consider the following example. Assuming that there is a constituent with parse tree $(p_2 \ (p_1 \ x \ y) \ z)$ (Figure 1), and $\mathbf{x}, \mathbf{y}, \mathbf{z} \in \mathbb{R}^{n \times 1}$ are respectively the meanings of the three words x, y and z. We will use a neural network which contains a weight matrix $\mathbf{W}_1 \in \mathbb{R}^{n \times n}$ for left children and a weight matrix $\mathbf{W}_2 \in \mathbb{R}^{n \times n}$ for right children to compute parents in a bottom up manner. Firstly, we use this network to compute p_1's meaning

$$\mathbf{p}_1 = f(\mathbf{W}_1 \mathbf{x} + \mathbf{W}_2 \mathbf{y} + \mathbf{b}) \qquad (2)$$

where $\mathbf{b}$ is a bias vector, f is an activation function (e.g. *tanh* or *logistic*). Then, we use the same network to compute p_2's meaning

$$\mathbf{p}_2 = f(\mathbf{W}_1 \mathbf{p}_1 + \mathbf{W}_2 \mathbf{z} + \mathbf{b}) \qquad (3)$$

This process is continued until we reach the root node. This network is trained by a gradient-based optimization method (e.g., gradient descent) where the gradient over parameters is efficiently computed thanks to the backpropagation through structure (Goller and Kuchler, 1996). Using this architecture (and its extensions), Socher and colleagues successfully reach state-of-the-art results in syntactic parsing (Socher et al., 2013a) and sentiment analysis (Socher et al., 2013b).


Figure 1: Recursive Neural Network (RNN).

In order to use this architecture in an unsupervised learning manner, Socher et al., 2011b) replace the neural network in RNN by an autoencoder (and hence the new architecture is called Recursive Autoencoder - RAE), which is a feedforward neural network trained by forcing output equal to input (see Figure 2). Training a RAE is therefore to minmize the sum of reconstruction errors (i.e., $\|[\mathbf{x}'; \mathbf{y}'] - [\mathbf{x}; \mathbf{y}]\|^2$) at all internall nodes.


Figure 2: Recursive Autoencoder (RAE).

4.2 Inside Outside Recursive Neural Network (IORNN)

None of the above architectures, RAE or RNN, compute contex representations. However, they give us a hint to do that. In this section, we will answer the second question "How to construct context representations?" by a new neural network architecture, namely Inside Outside Recursive Neural Network (IORNN). We also present this architecture by using the same example of constituent and parse tree $(p_2 \ (p_1 \ x \ y) \ z)$ (see Figure 3).

Each node u is assigned two vectors $\mathbf{o}_u$ and $\mathbf{i}_u$. The first one, called *outer meaning*, denotes the

Figure 3: Inside-Outside Recursive Network (IORNN). Black rectangles correspond to inner meanings, white rectangles correspond to outer meanings.

meaning of the context; the second one, called *inner meaning*, denotes the meaning of the phrase that the node covers.

Word embeddings (e.g., $\mathbf{i}_x$) Similar to Socher et al., 2010; Collobert et al., 2011), given a string of binary representations of words $(a, b, ..., w)$ (i.e., all of the entries of w are zero except the one corresponding to the index of the word in the vocabulary), we first compute a string of vectors $(\mathbf{i}_a, ..., \mathbf{i}_w)$ representing inner meanings of those words by using a look-up table (i.e., word embeddings) $\mathbf{L} \in \mathbb{R}^{n \times |V|}$, where $|V|$ is the size of the vocabulary and n is the dimensionality of the vectors. This look-up table $\mathbf{L}$ could be seen as a storage of word meanings where each column is a vector representation of a word. Hence,

$$\mathbf{i}_w = \mathbf{L}w \in \mathbb{R}^{n \times 1} \qquad (4)$$

Computing inner meaning The inner meaning of a non-terminal node, say p_1, is given by

$$\mathbf{i}_{p_1} = f(\mathbf{W}_1^i \mathbf{i}_x + \mathbf{W}_2^i \mathbf{i}_y + \mathbf{b}^i) \qquad (5)$$

where $\mathbf{W}_1^i$, $\mathbf{W}_2^i$ are $n \times n$ real matrices, $\mathbf{b}^i$ is a bias vector, and f is an activation function, e.g. $tanh$. Intuitively, the inner meaning of a parent node is the function of the inner meanings of its children. This is similar to RNN.

Computing outer meaning Because we process sentences individually, there is no information about discourse context. Therefore, the outer meaning of the root node, $\mathbf{o}_{root}$, is set to $\mathbf{o}_\emptyset$, which is randomly initialized and then learnt later. To a

node which is not the root, say p_1, the outer meaning is given by

$$\mathbf{o}_{p_1} = g(\mathbf{W}_1^o \mathbf{o}_{p_2} + \mathbf{W}_2^o \mathbf{i}_z + \mathbf{b}^o) \qquad (6)$$

where $\mathbf{W}_1^o$, $\mathbf{W}_2^o$ are $n \times n$ real matrices, $\mathbf{b}^o$ is a bias vector, and g is an activation function. Informally speaking, the outer meaning of a node (i.e., the meaning of its context) is the function of the outer meaning of its parent and the inner meaning of its sister.

The reader could recognize the similarity between Equation 5, 6 and the inside, outside probabilities given a parse tree. That is why we name the architecture Inside-Outside Recursive Neural Network.

4.3 Training IORNN

This section is to answer the final question "How to use words and their contexts to learn compositionality functions?".

According to Hypothesis 1, there must be a strong correlation between $\mathbf{o}_w$ and $\mathbf{i}_w$ where w is any word in a given sentence. The simplest way to train the network is to force $\mathbf{o}_{w_j} = \mathbf{i}_{w_j}$; hence, learning is to minimize the following loss function

$$J(\theta) = \sum_{s \in D} \sum_{w \in s} \|\mathbf{o}_w - \mathbf{i}_w\| \qquad (7)$$

where D is a set of training sentences and θ are the network parameters. However, that could be problematic because the meaning of context is not necessary the meaning of the target word.

Here, based on the observation that the meaning of context sets constraints on selecting a word to fill in the blank, one could suggest put a *softmax* neuron unit on the top of each $\mathbf{o}_w$ in order to compute the probability $P(x|\mathbf{o}_w)$. Unfortunately, as pointed out by Collobert et al., 2011), it might not work.

Using the same method proposed by Collobert et al., 2011), we train the network such that it gives to the correct target word a score higher than scores given to other words by a margin of 1. The score $s(x, \mathbf{o}_w)$ given to a candidate word x for a specific context $\mathbf{o}_w$ is computed by

$$u(x, \mathbf{o}_w) = f(\mathbf{W}_1^u \mathbf{o}_w + \mathbf{W}_2^u \mathbf{i}_x + \mathbf{b}^u) \qquad (8)$$

$$s(x, \mathbf{o}_w) = \mathbf{W}^s u(x, \mathbf{o}_w) + \mathbf{b}^s \qquad (9)$$

where $\mathbf{W}_1^u$, $\mathbf{W}_2^u$ are $n \times k$ real matrices, $\mathbf{W}^s$ is a $k \times 1$ matrix, and $\mathbf{b}^u$, $\mathbf{b}^s$ are bias vectors. (We

fix $k = 2n$.) Now, the objective function is the ranking criterion with respect to θ

$$J(\theta) = \sum_{s \in D} \sum_{w \in s} \sum_{x \in V} \max\{0, 1 - s(w, \mathbf{o}_w) + s(x, \mathbf{o}_w)\}$$
(10)

To minimize the above objective function, we randomly pick up a word in the vocabulary as a corrupt example, then compute the gradient by the backpropagation through structure (Goller and Kuchler, 1996). Following Socher et al., 2013b), we use AdaGrad (Duchi et al., 2011) to update the parameters.

5 Experiments

In order to examine how IORNN performs for both compositional semantics learning and meaning in context, we evaluated it in two tasks: phrase similarity and word meaning in context. Because, to our knowledge, there are no frameworks that tackle both problems, we use vector addition and pair-wise multiplication as our baselines. Although these methods are simple, choosing them as baselines are reasonable since (1) Blacoe and Lapata, 2012) show that they perform better than RAE in the phrase similarity task, (2) they are used in applications requiring compositional semantics, e.g. in the work of Šarić et al., 2012), and (3) vector addition is used as a method to compute contextualized vectors (Thater et al., 2011).

In the all experiments, we implemented IORNN in Torch-Lua (Collobert et al., 2012). We initialized the network with the Collobert & Weston (C&W) 50-dim word embeddings[2] from Collobert et al., 2011). Then we trained it on a dataset containing 1.5M sentences from the BNC corpus (about one fourth of the whole corpus), which were parsed by the Berkeley parser (Petrov et al., 2006) and binarized.

5.1 Qualitative Evaluations

First of all, to show that IORNN is capable to use context to predict the meaning of a unknown word, we run the trained network on the WSJ section 22 and measured the average predicted rank of target/gold-standard words. (We did not use sentences from the BNC corpus because we also wanted to examine the generality of IORNN: "Does it work on different domains?") For each target word, we create a list of 20,000 candidate

words consisting of 19,999 words randomly selected from the vocabulary and the target word itself. Then, we use the scores given by Equation 9 to rank those candidates.

We found the average predicted rank of target words is 203.4 over 20,000 ($\simeq 1\%$), which means that the context meaning computed by IORNN tends to prefer the correct word. Looking into details (see Table 1), we discovered that IORNN tends to predict correctly word class (e.g., noun in examples 1 and 3, verb in example 4, adverb in example 2), word form (e.g., plural in example 1 and 3, bare verb in example 4). This is interesting because grammatical categories are totally ignored in both training and test phases. In addition, in some cases, high rank words seems to be logical in corresponding contexts (e.g., example 3).

5.2 Phrase Similarity

Phrase similarity is the task in which one is asked to compute the meanings of (short) phrases and measure their semantic similarities. Its goodness is measured by comparing its judgements with human judgments. In this experiment, we used the dataset[3] from Mitchell and Lapata, 2010) which contains 5832 human judgments on semantic similarity for noun-noun, verb-object, and adjective-noun phrases. There are 108 items; each contains a phrase pair and human ratings from 1 (very low similarity) to 7 (very high similarity) (see Table 2).

In this task, we use the cosine distance to measure the semantic similarity, i.e. $d(a, b) = \cos(\mathbf{i}_a, \mathbf{i}_b)$. Following Blacoe and Lapata, 2012), Hermann and Blunsom, 2013) and many others, we compute Spearman's correlation coefficient ρ between model scores and human judgments.

type	phrase 1	phrase 2	rating
v-obj	remember name	pass time	3
adj-n	dark eye	left arm	5
n-n	county council	town hall	4

Table 2: Items in the dataset from Mitchell and Lapata, 2010).

First of all, we focus on the results reported by Blacoe and Lapata, 2012). They claim that RAE performs worse than addition and pair-wise multiplication in all of their three settings. Here, we used their third setting, and, to be fair, we trained

ID	word	top 10 candidates (out of 20,000)
1	Institutional investors and bankers [...] were cautiously optimistic after the mild 1.8% decline in Tokyo stock *prices*.	standards, hours, projects, roads, members, duty, restrictions, *prices*, house, locations
2	That is *why* everybody was a little surprised by the storm of sell orders from small private investors [...]	at, over, when, since, *why*, why, one, was, a
3	That is why everybody was a little surprised by the storm of sell orders from small private *investors*, " said Norbert Braeuer [...]	men, companies, courts, camps, businesses, sports, jobs, parks, partners, air
4	[...] most investors wanted to see what would *happen* in New York before acting.	go, say, try, and, want, ',', work, invest, to, forget

Table 1: Words, their contexts, and top 10 candidates.

IORNN with their neural language model word embeddings[4]. The results[5] are given in Table 3. IORNN is the best for adj-n and v-obj, and second for noun-noun.

dim.	model	adj-n	n-n	v-obj
50	add.	0.28	**0.26**	0.24
50	mult.	0.26	0.22	0.18
100	RAE	0.20	0.18	0.14
50	IORNN	**0.30**	0.23	**0.28**

Table 3: Spearman's correlation coefficients of model predictions for the phrase similarity task with neural language model word embeddings from Blacoe and Lapata, 2012).

Hermann and Blunsom, 2013) extend RAE with the help of Combinatory Categorial Grammar (CCG). Their models, named Combinatory Categorial Autoencoder (CCAE), are similar to RAE but use different parameter sets for different grammatical rules and grammatical types. Thank to this semantic-related grammar, their models outperform RAE and score towards the upper end of the range of addition and pair-wise multiplication. Table 4 shows the comparison between IORNN and other methods which results are copied from the corresponding papers.

IORNN's performance lies in the range of CCAEs for adj-n and v-obj, but worse for noun-noun. However, it is worth emphasizing that IORNN uses one parameter set for all grammat-

model	adj-n	n-n	v-obj
Blacoe & Lapata			
a./m.	0.21 - 0.48	0.22 - 0.50	0.18 - 0.35
RAE	0.20 - 0.34	0.18 - 0.29	0.06 - 0.32
Hermann & Blunsom			
CCAEs	0.38 - 0.41	0.41 - 0.44	0.23 - 0.34
Our implementation			
add.	0.30	0.43	0.30
mult.	0.14	0.24	0.16
IORNN	0.38	0.36	0.32

Table 4: Spearman's correlation coefficients of model predictions for the phrase similarity task.

ical rules (similarly to RAE). From the difference of performance between RAE and CCAEs, we expect that extending IORNN in the same way (i.e., using CCG and different parameter sets for different grammatical rules and grammatical types) will lead to better performance.

5.3 Word Similarity in Context

Differing from the first task, this task focuses on word meaning in context: it examines how well a model can make use of context to disambiguate word senses. In this experiment, we use the Stanford Word Similarity in Context (SWSC) dataset from Huang et al., 2012) which contains 2003 word pairs, their sentential contexts, and human ratings from 0 to 10 (see Table 5).

For IORNN, we represent the meaning of a word in its sentential context by concatenate its inner and outer meanings, i.e. $\mathbf{m}_w = [\mathbf{i}_w; \mathbf{o}_w]$. For vector addition, we compute the context meaning by averaging the meaning vectors of 5 words on the left and 5 words on the right and then concate-

word 1	word 2	human ratings
Located downtown along the east *bank* of the Des Moines River, the plaza is available for parties, ...	This is the basis of all *money* laundering, ...	0.0, 0.0, 3.0, 10.0, 8.0, 0.0, 4.0, 0.0, 0.0, 0.0

Table 5: An example from the SWSC dataset.

nate it with the meaning vector of the target word.

Similarly to the first experiment, we also use the cosine distance to measure the semantic similarity and compute Spearman's correlation coefficient ρ between model scores and human judgments.

We compare IORNN and vector addition with HSMN-M AvgSimC proposed by Huang et al., 2012) and Pruned tf-idf-M AvgSimC proposed by Reisinger and Mooney, 2010). These two methods are multi-prototype approaches: the meaning of a word is represented by multiple vectors (i.e., prototypes). In order to extract multiple prototypes for a word, they compute a vector for each context that the word is in, then cluster those context vectors. In this way, a prototype corresponding to the centroid of a cluster represents a sense of the word. In order to compute the word meaning similarity with context, they use AvgSimC metric

$$\text{AvgSimC}(w, w') =$$
$$\frac{1}{k^2} \sum_{i=1}^{k} \sum_{j=1}^{k} p(c, w, i) p(c', w', j) d(\mu_i(w), \mu_j(w'))$$

where k is the number of prototypes of each word, $p(c, w, i)$ is the likelihood that word w is in its cluster i given context c, $\mu_i(w)$ is the i-th cluster centroid of w and $d(v, v')$ is a function computing similarity between two vectors.

Table 6 shows the comparison between the methods. It is not surprising to see HSMN-M AvgSimC and Pruned tf-idf-M AvgSimC perform best since disambiguating word sense is the key to success and these methods use different vectors to represent different senses of a word. However, IORNN, which represents the meaning of a word by a vector, performs comparably with Pruned tf-idf-M AvgSimC, and higher than the vector addition. It is worth noting the improvement from without using context (C&W) to using context meaning computed by IORNN (from 58.0 to 60.2), thus confirms the ability of IORNN in computing word meaning in context.

Model	$\rho \times 100$
C&W (w/o context)	58.0
Huang et al.	
HSMN-M AvgSimC	65.7
Pruned tf-idf-M AvgSimC	60.5
Our implementation	
add.	59.0
IORNN	60.3

Table 6: Spearman's correlation coefficients of model predictions on the SWCS dataset. C&W is the method to use the Collobert & Wetson word embeddings without taking context into account.

5.4 Summary

Now, we combine the experimental results presented above to compare IORNN with the two baselines, vector addition and vector pair-wise multiplication (see Table 7). IORNN outperforms the both baselines for adj-n, v-obj in phrase similarity task and in word similarity in context task, and worse than vector addition for noun-noun. These results show us that we can tackle the two problems, compositional semantics and meaning in context, with the unified framework IORNN.

Model	Phrase similarity			WSC
	adj-n	n-n	v-obj	
add.	0.30	**0.43**	0.30	59.0
mult.	0.14	0.24	0.16	-
IORNN	**0.38**	0.36	**0.32**	**60.3**

Table 7: Comparison of IORNN with vector addition and pair-wise multiplication in the two tasks.

6 Discussion

In this section, we will discuss two important issues. The first one is about the cognitive plausibility of IORNN. The second is about potential extensions for it.

6.1 Cognitive Plausibility

We found that it is cognitively plausible to represent context meaning separately from phrase/word meaning, which we have called outer meaning and inner meaning respectively. The key point here is what is called *variable binding* in connectionsim (Smolensky, 1990) where, in our case, outer meanings could be seen as slots and inner meanings as fillers (see Figure 4). Similar to hierarchical prediction network proposed by Borensztajn et al., 2009), if an outer meaning and an inner meaning are strongly correlated, a binding occurs and connects the neurons at the tree root node of the phrase to the neurons at the corresponding node of the tree of the context. This explains why a human being can use context meaning to predict the meaning of a unknown word, and why (s)he can select a word/phrase to fill in a blank in a incomplete sentence.


Figure 4: If there is a strong correlation between outer meaning and inner meaning, a dynamic binding occurs.

6.2 Potential Extensions / Future Work

Our new architecture IORNN has many potential extensions, some of which are promising to improve the performance in those two tasks presented above, other could lead to important applications.

At lexicon level In Subsection 5.3, we have seen multi-prototype approaches are promising for computing word meaning in context. Certainly, we can combine these approaches with our framework IORNN at the lexical level in order to disambiguate word senses.

At syntax level IORNN only uses parse trees without grammatical categories. However, Hermann and Blunsom, 2013) empirically show the important role of syntax in vector space models of compositional semantics. It turns out that it is also easy to extend IORNN in the same way, i.e. using CCG and different parameter sets for different grammatical rules and grammatical types. Thanks to some degree of similarity between IORNN and RAE, we expect that this extension helps IORNN improve its capacity in capturing compositionality.

At discourse level In this paper, we propose IORNN as an architecture processing individual sentences; therefore, the outer meaning at the root node is alway a null-context outer meaning vector (i.e., $\mathbf{o}_{root} = \mathbf{o}_{\emptyset}$). Beyond that, we can extend IORNN to make use of discourse context. Figure 5 illustrates how to connect inner and outer meanings of sentences in a discourse. Intuitively, the outer meaning of a sentence is the function of the inner meanings of its neighbours.


Figure 5: IORNN with discourse context. Black rectangles correspond to inner meanings, white rectangles correspond to outer meanings.

7 Conclusion

In this paper, we presented a unified framework, namely Inside Outside Recursive Neural Network, that tackles both compositional semantics and meaning in context. Similar to RNN, IORNN uses recursive network architecture to construct phrase representations, but different in that it also constructs context representations and is trained in a very different manner. The training, based on our new hypothesis, only needs contexts of words and thus avoids the sparsity of data, unlike the works of Baroni et al., 2012), Grefenstette et al., 2013).

References

Baroni, M., Bernardi, R., and Zamparelli, R. (2012). Frege in space: A program for compositional distributional semantics.

Blacoe, W. and Lapata, M. (2012). A comparison of vector-based representations for semantic composition. *RAE*, 70(68.29):6910.

Borensztajn, G., Zuidema, W., and Bod, R. (2009). The hierarchical prediction network: towards a neural theory of grammar acquisition. In *31th Annual Conference of the Cognitive Science Society*. Citeseer.

Brown, P. F., Desouza, P. V., Mercer, R. L., Pietra, V. J. D., and Lai, J. C. (1992). Class-based n-gram models of natural language. *Computational linguistics*, 18(4):467–479.

Collobert, R., Kavukcuoglu, K., and Farabet, C. (2012). Implementing neural networks efficiently. In *Neural Networks: Tricks of the Trade*, page 537557. Springer.

Collobert, R., Weston, J., Bottou, L., Karlen, M., Kavukcuoglu, K., and Kuksa, P. (2011). Natural language processing (almost) from scratch. *The Journal of Machine Learning Research*, 12:24932537.

Dhillon, P., Rodu, J., Foster, D., and Ungar, L. (2012). Two step cca: A new spectral method for estimating vector models of words. *arXiv preprint arXiv:1206.6403*.

Duchi, J., Hazan, E., and Singer, Y. (2011). Adaptive subgradient methods for online learning and stochastic optimization. *The Journal of Machine Learning Research*, 999999:2121–2159.

Garrette, D., Erk, K., and Mooney, R. (2012). A formal approach to linking logical form and vector-space lexical semantics. *Computing Meaning*, 4.

Goller, C. and Kuchler, A. (1996). Learning task-dependent distributed representations by backpropagation through structure. In *Neural Networks, 1996., IEEE International Conference on*, volume 1, page 347352.

Grefenstette, E., Dinu, G., Zhang, Y.-Z., Sadrzadeh, M., and Baroni, M. (2013). Multi-step regression learning for compositional distributional semantics. *arXiv preprint arXiv:1301.6939*.

Hermann, K. M. and Blunsom, P. (2013). The role of syntax in vector space models of compositional semantics. *Proceedings of ACL, Sofia, Bulgaria, August. Association for Computational Linguistics*.

Huang, E. H., Socher, R., Manning, C. D., and Ng, A. Y. (2012). Improving word representations via global context and multiple word prototypes. In *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics: Long Papers-Volume 1*, pages 873–882. Association for Computational Linguistics.

Landauer, T. K., Foltz, P. W., and Laham, D. (1998). An introduction to latent semantic analysis. *Discourse processes*, 25(2-3):259–284.

Le, P. and Zuidema, W. (2012). Learning compositional semantics for open domain semantic parsing. In *COLING*, pages 1535–1552.

Lenci, A. (2008). Distributional semantics in linguistic and cognitive research. *Italian journal of linguistics*, 20(1):131.

Mitchell, J. and Lapata, M. (2008). Vector-based models of semantic composition. *proceedings of ACL-08: HLT*, page 236244.

Mitchell, J. and Lapata, M. (2010). Composition in distributional models of semantics. *Cognitive Science*, 34(8):13881429.

Montague, R. (1970). English as a formal language. *Linguaggi nella societae nella tecnica*, pages 189–224.

Partee, B. (1995). Lexical semantics and compositionality. *An invitation to cognitive science: Language*, 1:311360.

Petrov, S., Barrett, L., Thibaux, R., and Klein, D. (2006). Learning accurate, compact, and interpretable tree annotation. In *Proceedings of the 21st International Conference on Computational Linguistics and the 44th annual meeting of the Association for Computational Linguistics*, pages 433–440. Association for Computational Linguistics.

Reisinger, J. and Mooney, R. J. (2010). Multi-prototype vector-space models of word meaning. In *Human Language Technologies: The 2010 Annual Conference of the North American Chapter of the Association for Computational Linguistics*, pages 109–117. Association for Computational Linguistics.

Šarić, F., Glavaš, G., Karan, M., Šnajder, J., and Bašić, B. D. (2012). Takelab: Systems for measuring semantic text similarity. In *Proceedings of the First Joint Conference on Lexical and Computational Semantics-Volume 1: Proceedings of the main conference and the shared task, and Volume 2: Proceedings of the Sixth International Workshop on Semantic Evaluation*, pages 441–448. Association for Computational Linguistics.

Smolensky, P. (1990). Tensor product variable binding and the representation of symbolic structures in connectionist systems. *Artificial intelligence*, 46(1):159–216.

Socher, R., Bauer, J., Manning, C. D., and Ng, A. Y. (2013a). Parsing with compositional vector grammars. In *Annual Meeting of the Association for Computational Linguistics (ACL)*.

Socher, R., Manning, C. D., and Ng, A. Y. (2010). Learning continuous phrase representations and syntactic parsing with recursive neural networks. In *Proceedings of the NIPS-2010 Deep Learning and Unsupervised Feature Learning Workshop*.

Socher, R., Pennington, J., Huang, E. H., Ng, A. Y., and Manning, C. D. (2011a). Semi-supervised recursive autoencoders for predicting sentiment distributions. In *Proceedings of the Conference on Empiri-*

cal Methods in Natural Language Processing, page 151161.

Socher, R., Pennington, J., Huang, E. H., Ng, A. Y., and Manning, C. D. (2011b). Semi-supervised recursive autoencoders for predicting sentiment distributions. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, pages 151–161. Association for Computational Linguistics.

Socher, R., Perelygin, A., Wu, J. Y., Chuang, J., Manning, C. D., Ng, A. Y., and Potts, C. (2013b). Recursive deep models for semantic compositionality over a sentiment treebank. EMNLP.

Thater, S., Fürstenau, H., and Pinkal, M. (2011). Word meaning in context: A simple and effective vector model. In *IJCNLP*, pages 1134–1143.