

Inside Outside Recursive Neural Network: A Unified Framework for Compositional Semantics and Meaning in Context

Phong Le, Willem Zuidema

1 Introduction

[...]

Firstly, we ask ourselves “Which evidence is strong enough to be used for learning compositional semantics?” To answer this question, we rely on the following observation: a human being can guess the meaning of an unknown word by making use of the meaning of its context. In other words, he computes (in his brain) the meaning of the context and then use it to predict the meaning of the target unknown word (by setting some constraints to narrow down a list of possible meanings). Hence, if he correctly predicts the meaning of the unknown word, we can, at some level of belief, say that he coherends the meaning of the context. This idea is then encapsured in the following hypothesis

Hypothesis 1: The agreement between words and contexts provides evidence for unsupervised compositional semantics learning.

Then, there are three questions need to be answered in order to implement the hypothesis

1. How to construct phrasal semantics?
2. How to construct contextual semantics?
3. How to use contextual semantics and lexical semantics to learn compositionality functions?

In the next section, we answer these questions.

2 Recursive Neural Network (RNN)

(Socher et al., 2010) answer the first question “How to construct phrasal semantics?” by Recursive Neural Network (RNN) architecture, which is used to compute continuous phrase representations. In order to see how RNN works, let’s consider the following example. Assuming that there is a constituent with parse tree $(p_2 (p_1 x y) z)$ as




Figure 1: Example of RNN.

in Figure 1. We will use a neural network, which contains a weight matrix \mathbf{W}_1 for left children and a weight matrix \mathbf{W}_2 for right children, to compute parents in a bottom up manner. Firstly, we use this network to compute p_1 based on its children x and y

$$\mathbf{p}_1 = f(\mathbf{W}_1\mathbf{x} + \mathbf{W}_2\mathbf{y} + \mathbf{b}) \quad (1)$$

where \mathbf{b} is a bias, f is an activation function (e.g. *tanh* or *logistic*). Then, we use the same network to compute p_2 based on its children p_1 and z

$$\mathbf{p}_2 = f(\mathbf{W}_1\mathbf{p}_1 + \mathbf{W}_2\mathbf{z} + \mathbf{b}) \quad (2)$$

This process is continued until we reach the root node. This network is trained by a gradient-based optimization method (e.g., gradient descent) where the gradient over parameters is efficiently computed thanks to the backpropagation through structure (Goller and Kuchler, 1996). Using this architecture (and its extensions), Socher and colleagues successfully reach state-of-the-art results in syntactic parsing (Socher et al., 2013a) and sentiment analysis (Socher et al., 2013b).

In order to use this architecture in an unsupervised learning manner, (Socher et al., 2011) replace neural network in RNN by autoencoder (and hence the new architecture is called Recursive Autoencoder - RAE), which is a feedforward neu-




Figure 2: Example of RAE.

ral network trained by forcing output equal to input (see Figure 2). Training a RAE is therefore to minimize the sum of reconstruction errors (i.e., $\|[\mathbf{x}'; \mathbf{y}'] - [\mathbf{x}; \mathbf{y}]\|^2$) at all internal nodes.

3 Inside Outside Recursive Neural Network (IORNN)

None of the above architectures, RAE or RNN, compute contextual semantics. However, they give us a hint to do that. In this section, we will answer the second question “How to construct contextual semantics?” by a new neural network architecture, namely Inside Outside Recursive Neural Network (IORNN). We also present this architecture by using an example of a constituent and parse tree ($p_2 (p_1 x y) z$) (see Figure 3).




Figure 3: Inside-Outside Recursive Neural Network. Black rectangles correspond to inner meanings, white rectangles correspond to outer meanings.

Each node u is assigned two vectors \mathbf{o}_u and \mathbf{i}_u . The first one, called *outer meaning*, denotes the meaning of the context; the second one, called *in-*

ner meaning, denotes the meaning of the phrase that the node covers.

Word embeddings (e.g., \mathbf{i}_x) Similar to (Socher et al., 2010), and (Collobert et al., 2011), given a string of binary representations of words (a, b, \dots, w) (i.e., all of the entries of w are zero except the one corresponding to the index of the word in the dictionary), we first compute a string of vectors ($\mathbf{i}_a, \dots, \mathbf{i}_w$) representing inner meanings of those words by using a look-up table (i.e., word embeddings) $\mathbf{L} \in \mathbb{R}^{n \times |V|}$, where $|V|$ is the size of the vocabulary and n is the dimensionality of the vectors. This look-up table \mathbf{L} could be seen as a storage of lexical semantics where each column is a vector representation of a word. Hence,

$$\mathbf{i}_w = \mathbf{L}w \in \mathbb{R}^n \quad (3)$$

Computing inner meaning The inner meaning of a non-terminal node, say p_1 , is given by

$$\mathbf{i}_{p_1} = f(\mathbf{W}_1^i \mathbf{i}_x + \mathbf{W}_2^i \mathbf{i}_y + \mathbf{b}^i) \quad (4)$$

where $\mathbf{W}_1^i, \mathbf{W}_2^i$ are $n \times n$ real matrices, \mathbf{b}^i is a bias vector, and $f(\cdot)$ is an activation function, e.g. \tanh function. Intuitively, the inner meaning of a parent node is the function of the inner meanings of its children. This is similar to what (Socher et al., 2010) call recursive neural network.

Computing outer meaning The outer meaning of the root node, \mathbf{o}_{root} , is initially set randomly, and then learnt later. To a node which is not the root, say p_1 , the outer meaning is given by

$$\mathbf{o}_{p_1} = g(\mathbf{W}_1^o \mathbf{o}_{p_2} + \mathbf{W}_2^o \mathbf{i}_z + \mathbf{b}^o) \quad (5)$$

where $\mathbf{W}_1^o, \mathbf{W}_2^o$ are $n \times n$ real matrices, \mathbf{b}^o is a bias vector, and $g(\cdot)$ is an activation function, e.g. \tanh function. Informally speaking, the outer meaning of a node (i.e., the meaning of its context) is the function of the outer meaning of its parent and the inner meaning of its sister.

The reader, if familiar with syntactic parsing, could recognizes the similarity between Equation 4, 5 and the inside, outside probabilities given a parse tree. Therefore, we name the architecture Inside-Outside Recursive Neural Network.

4 Training IORNN

This section is to answer the final question “How to use contextual semantics and lexical semantics to learn compositionality functions?”.

According to Hypothesis 1, there must be a strong correlation between \mathbf{o}_w and \mathbf{i}_w where w is any word in a given sentence. The simplest way to train the network is to force $\mathbf{o}_{w_j} = \mathbf{i}_{w_j}$; hence, learning is to minimize the following loss function

$$J(\theta) = \sum_{s \in D} \sum_{w \in s} \|\mathbf{o}_w - \mathbf{i}_w\| \quad (6)$$

where D is a set of training sentences and θ are the network parameters. However, that could be problematic because the meaning of context is not necessary the meaning of the target word.

Here, based on the observation that the meaning of context sets constraints on selecting a word to fill in the blank, one could suggest put a *softmax* neuron unit on the top of each \mathbf{o}_w in order to compute the probability $P(x|\mathbf{o}_w)$. Unfortunately, as pointed out by (Collobert et al., 2011), it might not work.

Using the same method proposed by (Collobert et al., 2011), we train the network such that it gives a higher score to the correct target word rather than to incorrect ones. The score $s(x, \mathbf{o}_w)$ given to a candidate word x for a specific context \mathbf{o}_w is computed by

$$u(x, \mathbf{o}_w) = f(\mathbf{W}_u^o \mathbf{o}_w + \mathbf{W}_u^i \mathbf{i}_x + \mathbf{b}_u) \quad (7)$$

$$s(x, \mathbf{o}_w) = \mathbf{W}_s u(x, \mathbf{o}_w) + \mathbf{b}_s \quad (8)$$

where $\mathbf{W}_u^o, \mathbf{W}_u^i$ are $n \times k$ real matrices, \mathbf{W}_s is a $k \times 1$ matrix, and $\mathbf{b}_u, \mathbf{b}_s$ are bias vectors. (We fix $k = 2n$.) Now, the objective function is the ranking criterion with respect to θ

$$J(\theta) = \sum_{s \in D} \sum_{w \in s} \sum_{x \in V} \max\{0, 1 - s(w, \mathbf{o}_w) + s(x, \mathbf{o}_w)\} \quad (9)$$

To minimize the above objective function, we randomly pick up a word in the dictionary as a corrupt example, compute the gradient, and update the parameters by a gradient descent method. Thanks to the backpropagation through structure (Goller and Kuchler, 1996), the gradient is efficiently computed. Following (Socher et al., 2013b), we use AdaGrad (Duchi et al., 2011) to update the parameters.

References

Collobert, R., Weston, J., Bottou, L., Karlen, M., Kavukcuoglu, K., and Kuksa, P. (2011). Natural language processing (almost) from scratch. *The Journal of Machine Learning Research*, 12:24932537.

Duchi, J., Hazan, E., and Singer, Y. (2011). Adaptive subgradient methods for online learning and stochastic optimization. *The Journal of Machine Learning Research*, 999999:2121–2159.

Goller, C. and Kuchler, A. (1996). Learning task-dependent distributed representations by backpropagation through structure. In *Neural Networks, 1996., IEEE International Conference on*, volume 1, page 347352.

Socher, R., Bauer, J., Manning, C. D., and Ng, A. Y. (2013a). Parsing with compositional vector grammars. In *Annual Meeting of the Association for Computational Linguistics (ACL)*.

Socher, R., Manning, C. D., and Ng, A. Y. (2010). Learning continuous phrase representations and syntactic parsing with recursive neural networks. In *Proceedings of the NIPS-2010 Deep Learning and Unsupervised Feature Learning Workshop*.

Socher, R., Pennington, J., Huang, E. H., Ng, A. Y., and Manning, C. D. (2011). Semi-supervised recursive autoencoders for predicting sentiment distributions. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, pages 151–161. Association for Computational Linguistics.

Socher, R., Perelygin, A., Wu, J. Y., Chuang, J., Manning, C. D., Ng, A. Y., and Potts, C. (2013b). Recursive deep models for semantic compositionality over a sentiment treebank. EMNLP.